

CMPSC 442: Homework 5 [100 points]

Release Date Monday, March 23, 2020, 12:00 am

Due Date Friday, April 10, 2020, 11:59 pm

TO SUBMIT HOMEWORK

To submit homework for a given homework assignment:

1. You ***must*** download the homework template file from Canvas, located in Files/Homework Templates and Pdfs, and modify this file to complete your homework. Each template file is a python file that will give you a headstart in creating your homework python script. For a given homework number N, the template file name is homeworkN-cmpsc442.py. The template for homework #5 is homework5-cmpsc442.py. **IF YOU DO NOT USE THE CORRECT TEMPLATE FILE, YOUR HOMEWORK CANNOT BE GRADED AND YOU WILL RECEIVE A ZERO.** There is also a text file in the Files folder that you will need, called:
 - frankenstein.txt
2. You ***must*** rename the file by replacing the file root using your PSU id that consists of your initials followed by digits. This is the same as the part of your PSU email that precedes the "@" sign. For example, your instructor's email is rjp49@cse.psu.edu, and her PSU id is rjp49. Your homework files for every assignment will have the same name, e.g., rjp49.py. **IF YOU DO NOT RENAME YOUR HOMEWORK FILE CORRECTLY, IT WILL NOT BE GRADED AND YOU WILL RECEIVE A ZERO.** Do not be alarmed if you upload a revision, and it is renamed to include a numeric index, e.g., rjp49-1.py or rjp49-2.py. We can handle this automatic renaming.
3. You ***must*** upload your homework to the assignments area in Canvas by 11:59 pm on the due date. You will have two opportunities (NO MORE) to submit up to two days late. **IF YOU DO NOT UPLOAD YOUR HOMEWORK TO THE ASSIGNMENT FOLDER BY THE DUE DATE (OR THE TWO-DAY GRACE PERIOD IN SOME CASES), IT CANNOT BE GRADED AND YOU WILL RECEIVE A ZERO.**

Instructions

In this assignment, you will gain experience working with Markov models on text.

A skeleton file homework5-cmpsc442.py containing empty definitions for each question has been provided. Since portions of this assignment will be graded automatically, none of the names or function signatures in this file should be modified. However, you are free to introduce additional variables or functions if needed.

You may import definitions from any standard Python library, and are encouraged to do so in case you find yourself reinventing the wheel.

You will find that in addition to a problem specification, most programming questions also include a pair of examples from the Python interpreter. These are meant to illustrate **typical use cases** to clarify the assignment, and are **not comprehensive test suites**. In addition to performing your own testing, you are strongly encouraged to verify that your code gives the expected output for these examples before submitting.

You are strongly encouraged to follow the Python style guidelines set forth in [PEP 8](#), which was written in part by the creator of Python. However, your code will not be graded for style.

1. Markov Models [95 points]

In this section, you will build a simple language model that can be used to generate random text resembling a source document. Your use of external code should be limited to built-in Python modules, which excludes, for example, NumPy and NLTK.

1. **[10 points]** Write a simple tokenization function `tokenize(text)` which takes as input a string of text and returns a list of tokens derived from that text. Here, we define a token to be a contiguous sequence of non-whitespace characters, with the exception that any punctuation mark should be treated as an individual token. *Hint: Use the built-in constant `string.punctuation`, found in the `string` module.*

```
>>> tokenize(" This is an example. ")
['This', 'is', 'an', 'example', '.']
```

```
>>> tokenize("'Medium-rare,' she said.")
['', 'Medium', '-', 'rare', ',', '', 'she', 'said', '.']
```

2. **[10 points]** Write a function `ngrams(n, tokens)` that produces a list of all n -grams of the specified size from the input token list. Each n -gram should consist of a 2-element tuple (context, token), where the context is itself an $(n-1)$ -element tuple comprised of the $n-1$ words preceding the current token. The sentence should be padded with $n-1$ "`<START>`" tokens at the beginning and a single "`<END>`" token at the end. If $n = 1$, all contexts should be empty tuples. You may assume that $n \geq 1$.

```
>>> ngrams(1, ["a", "b", "c"])
[(), 'a'), ((), 'b'), ((), 'c'),
 ((), '<END>')]
>>> ngrams(2, ["a", "b", "c"])
[('<START>',), 'a'), (('a',), 'b'),
 (('b',), 'c'), (('c',), '<END>')]
```

```
>>> ngrams(3, ["a", "b", "c"])
[('<START>', '<START>'), 'a'),
 (('<START>', 'a'), 'b'),
 (('a', 'b'), 'c'),
 (('b', 'c'), '<END>')]
```

3. **[15 points]** In the `NgramModel` class, write an initialization method `__init__(self, n)` which stores the order n of the model and initializes any necessary internal variables. Then write a method `update(self, sentence)` which computes the n -grams for the input sentence and updates the internal counts. Lastly, write a method `prob(self, context, token)` which accepts an $(n-1)$ -tuple representing a context and a token, and returns the probability of that token occurring, given the preceding context.

```
>>> m = NgramModel(1)
>>> m.update("a b c d")
>>> m.update("a b a b")
>>> m.prob(), "a"
0.3
>>> m.prob(), "c"
0.1
>>> m.prob(), "<END>"
```

```
>>> m = NgramModel(2)
>>> m.update("a b c d")
>>> m.update("a b a b")
>>> m.prob(("<START>",), "a")
1.0
>>> m.prob(("b",), "c")
0.3333333333333333
>>> m.prob(("a",), "x")
```

0.2

0.0

4. **[20 points]** In the `NgramModel` class, write a method `random_token(self, context)` which returns a random token according to the probability distribution determined by the given context. Specifically, let $T = \langle t_1, t_2, \dots, t_n \rangle$ be the set of tokens which can occur in the given context, sorted according to Python's natural lexicographic ordering, and let $0 \leq r < 1$ be a random number between 0 and 1. Your method should return the token t_i such that

$$\sum_{j=1}^{i-1} P(t_j \mid \text{context}) \leq r < \sum_{j=1}^i P(t_j \mid \text{context}).$$

You should use a single call to the `random.random()` function to generate r .

```
>>>m = NgramModel(1)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>random.seed(1)
>>>[m.random_token(())
    for i in range(25)]
['<END>', 'c', 'b', 'a', 'a', 'a', 'b',
 'b', '<END>', '<END>', 'c', 'a', 'b',
 '<END>', 'a', 'b', 'a', 'd', 'd',
 '<END>', '<END>', 'b', 'd', 'a', 'a']
```

```
>>>m = NgramModel(2)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>random.seed(2)
>>>[m.random_token("<START>"),)
    for i in range(6)]
['a', 'a', 'a', 'a', 'a', 'a']
>>>[m.random_token(("b"),)
    for i in range(6)]
['c', '<END>', 'a', 'a', 'a', '<END>']
```

5. **[20 points]** In the `NgramModel` class, write a method `random_text(self, token_count)` which returns a string of space-separated tokens chosen at random using the `random_token(self, context)` method. Your starting context should always be the $(n-1)$ -tuple $(\text{"<START>"}, \dots, \text{"<START>"})$, and the context should be updated as tokens are generated. If $n = 1$, your context should always be the empty tuple. Whenever the special token `"<END>"` is encountered, you should reset the context to the starting context.

```
>>>m = NgramModel(1)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>random.seed(1)
>>>m.random_text(13)
'<END> c b a a a b b <END> <END> c a b'
```

```
>>>m = NgramModel(2)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>random.seed(2)
>>>m.random_text(15)
'a b <END> a b c d <END> a b a b a b c'
```

6. **[10 points]** Write a function `create_ngram_model(n, path)` which loads the text at the given path and creates an n -gram model from the resulting data. Each line in the file should be treated as a separate sentence.

```
# No random seeds, so your results may vary
```

```
>>>m = create_ngram_model(1, "frankenstein.txt"); m.random_text(15)
'beat astonishment brought his for how , door <END> his . pertinacity to I felt'
>>>m = create_ngram_model(2, "frankenstein.txt"); m.random_text(15)
'As the great was extreme during the end of being . <END> Fortunately the sun'
>>>m = create_ngram_model(3, "frankenstein.txt"); m.random_text(15)
'I had so long inhabited . <END> You were thrown , by returning with greater'
>>>m = create_ngram_model(4, "frankenstein.txt"); m.random_text(15)
'We were soon joined by Elizabeth . <END> At these moments I wept bitterly and'
```

7. **[10 points]** Suppose we define the perplexity of a sequence of m tokens $\langle w_1, w_2, \dots, w_m \rangle$ to be

$$\sqrt[m]{\frac{1}{P(w_1, w_2, \dots, w_m)}}.$$

For example, in the case of a bigram model under the framework used in the rest of the assignment, we would generate the bigrams $\langle (w_o = \langle \text{START} \rangle, w_1), (w_1, w_2), \dots, (w_{m-1}, w_m), (w_m, w_{m+1} = \langle \text{END} \rangle) \rangle$, and would then compute the perplexity as

$$\sqrt[m+1]{\prod_{i=1}^{m+1} \frac{1}{P(w_i | w_{i-1})}}.$$

Intuitively, the lower the perplexity, the better the input sequence is explained by the model. Higher values indicate the input was "perplexing" from the model's point of view, hence the term perplexity.

In the `NgramModel` class, write a method `perplexity(self, sentence)` which computes the n -grams for the input sentence and returns their perplexity under the current model. *Hint: Consider performing an intermediate computation in log-space and re-exponentiating at the end, so as to avoid numerical overflow.*

```
>>>m = NgramModel(1)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>m.perplexity("a b")
3.815714141844439
```

```
>>>m = NgramModel(2)
>>>m.update("a b c d")
>>>m.update("a b a b")
>>>m.perplexity("a b")
1.4422495703074083
```

2. Feedback [5 points]

1. **[1 point]** Approximately how long did you spend on this assignment?
2. **[2 points]** Which aspects of this assignment did you find most challenging? Were there any significant stumbling blocks?
3. **[2 points]** Which aspects of this assignment did you like? Is there anything you would have changed?